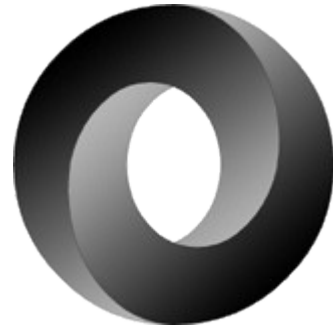


NoSQL as Not Only SQL

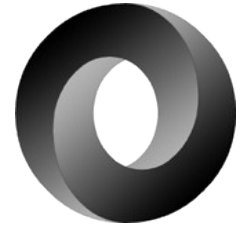


PostgreSQL

Backend Berlin
Dezember 2015
Stefanie Janine Stölting
[@sjstoelting](https://twitter.com/sjstoelting)



JSON Datatypes



JSON

Available since 9.2

BSON

Available as extension on GitHub since 2013

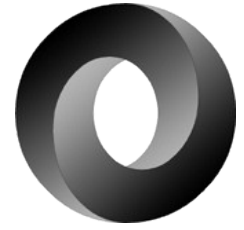
JSONB

Available since 9.4

Compressed JSON



ACID



Atomicity, Consistency, Isolation, Durability is a set of properties that guarantee that database transactions are processed reliably. ¹

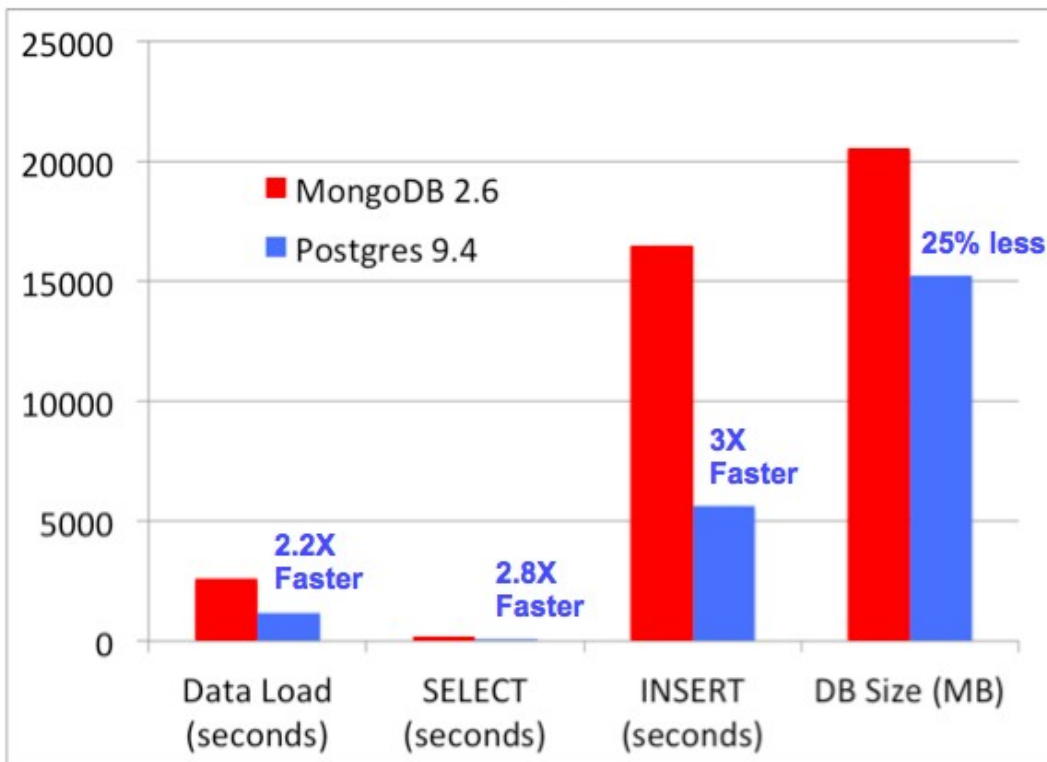
¹ See <https://en.wikipedia.org/wiki/ACID>



Performance



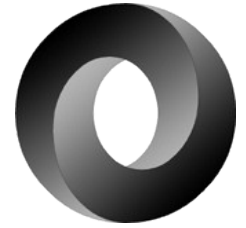
MongoDB 2.6 vs PostgreSQL 9.4 Performance



Test done by
[EnterpriseDB](#),
see the [article](#) by
[Marc Linster](#)



JSON Functions



`row_to_json({row})`

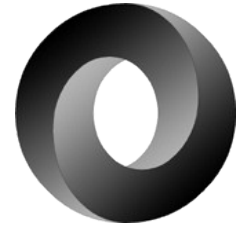
Returns the row as JSON

`array_to_json({array})`

Returns the array as JSON



JSON Operators



Array element

->{int}

Array element by name

->{text}

Object element

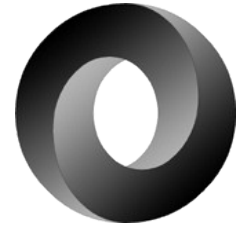
->> {text}

Value at path

#> {text}



Index on JSON



Index JSONB content for faster access

GIN index overall

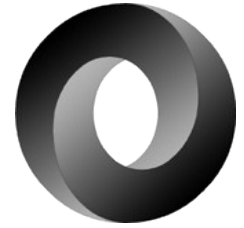
```
CREATE INDEX idx_1 ON jsonb.actor USING  
GIN (jsondata);
```

Even unique B-Tree indexes are possible

```
CREATE UNIQUE INDEX actor_id_2 ON  
jsonb.actor((CAST(jsondata → 'actor_id'::int)));
```



New JSON functions



With the next version, 9.5, there will be new functions available, for example:

`jsonb_pretty`

`jsonb_set`

If you can't wait using them, you might be interested in the [jsonb_x](#) extension available at PGXN



Data sources



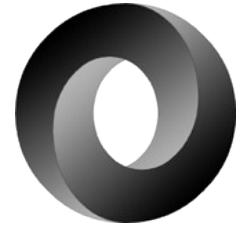
The Chinook database is available at chinookdatabase.codeplex.com

Amazon book reviews of 1998 are available at

examples.citusdata.com/customer_reviews_nested_1998.json.gz



CTE



Common Table Expressions will be used in examples

Example:

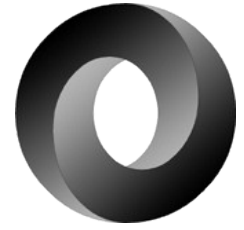
```
WITH RECURSIVE t(n) AS (  
    VALUES (1)  
    UNION ALL  
    SELECT n+1 FROM t WHERE n < 100  
)  
SELECT sum(n), min(n), max(n) FROM t;
```

Result:

	sum bigint	min integer	max integer
1	5050	1	100



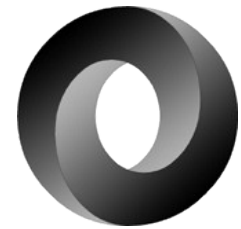
Live Examples



Let's see, how it does work.



Live with Chinook data



-- Step 1: Albums with tracks as JSON

```
WITH albums AS
(
    SELECT a."ArtistId" AS artist_id
        , a."Title" AS album_title
        , array_agg(t."Name") AS album_tracks
    FROM "Album" AS a
    INNER JOIN "Track" AS t
        ON a."AlbumId" = t."AlbumId"
    GROUP BY a."ArtistId"
        , a."Title"
)
SELECT row_to_json(albums) AS album_tracks
FROM albums
;
```


	album_tracks
1	{"artist_id":133,"album_title":"In Step","album_tracks":["The House Is Rockin'","Crossfire","Tightrope","Let Me Love You Baby","Leave My Girl Alone","T
2	{"artist_id":83,"album_title":"Deixa Entrar","album_tracks":["Deixa Entrar","Falamansa Song","Xote Dos Milagres","Zeca Violeiro","Avisa","Principiando/I
3	{"artist_id":15,"album_title":"The Best Of Buddy Guy - The Millenium Collection","album_tracks":["First Time I Met The Blues","Let Me Love You Baby","
4	{"artist_id":78,"album_title":"Vault: Def Leppard's Greatest Hits","album_tracks":["Pour Some Sugar On Me","Photograph","Love Bites","Let's Get Rock
5	{"artist_id":205,"album_title":"Carry On","album_tracks":["No Such Thing","Poison Eye","Arms Around Your Love","Safe and Sound","She'll Never Be You
6	{"artist_id":58,"album_title":"Come Taste The Band","album_tracks":["Comin' Home","Lady Luck","Gettin' Tighter","Dealer","I Need Love","Drifter","Lov
7	{"artist_id":110,"album_title":"Nevermind","album_tracks":["Smells Like Teen Spirit","In Bloom","Come As You Are","Breed","Lithium","Polly","Territoria

200 row(s) fetched - 14ms

Grid



Live with Chinook data



-- Step 2 Albums including tracks with artists

```
WITH albums AS
(
    SELECT a."ArtistId" AS artist_id
        , a."Title" AS album_title
        , array_agg(t."Name") AS album_tracks
    FROM "Album" AS a
        INNER JOIN "Track" AS t
            ON a."AlbumId" = t."AlbumId"
    GROUP BY a."ArtistId"
        , a."Title"
)
, js_albums AS
(
    SELECT row_to_json(albums) AS album_tracks
    FROM albums
)
SELECT a."Name" AS artist
    , al.album_tracks AS albums_tracks
FROM sqlite_artist AS a
    INNER JOIN js_albums AS al
        ON a."ArtistId" = CAST(al.album_tracks->>'artist_id' AS INT)
;
```


	artist	albums_tracks
1	AC/DC	{"artist_id":1,"album_title":"Let There Be Rock","album_tracks":["Go Down","Dog Eat Dog","Let There Be Rock","Bad Boy Boogie",
2	AC/DC	{"artist_id":1,"album_title":"For Those About To Rock We Salute You","album_tracks":["For Those About To Rock (We Salute You)
3	Accept	{"artist_id":2,"album_title":"Restless and Wild","album_tracks":["Fast As a Shark","Restless and Wild","Princess of the Dawn"]}
4	Accept	{"artist_id":2,"album_title":"Balls to the Wall","album_tracks":["Balls to the Wall"]}
5	Aerosmith	{"artist_id":3,"album_title":"Big Ones","album_tracks":["Walk On Water","Love In An Elevator","Rag Doll","What It Takes","Dude (
6	Alanis Morissette	{"artist_id":4,"album_title":"Jagged Little Pill","album_tracks":["All I Really Want","You Oughta Know","Perfect","Hand In My Pock
7	Alice In Chains	{"artist_id":5,"album_title":"Facelift","album_tracks":["We Die Young","Man In The Box","Sea Of Sorrow","Bleed The Freak","I Can
8	Apocalyptica	{"artist_id":7,"album_title":"Plays Metallica By Four Cellos","album_tracks":["Enter Sandman","Master Of Puppets","Harvester Of

200 row(s) fetched - 48ms

Grid



Live with Chinook data

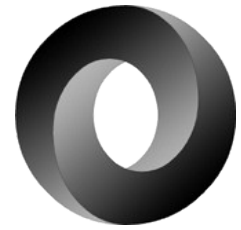


-- Step 3 Return one row for an artist with all albums as VIEW

```
CREATE OR REPLACE VIEW v_artist_data AS
WITH albums AS
(
    SELECT a."ArtistId" AS artist_id
        , a."Title" AS album_title
        , array_agg(t."Name") AS album_tracks
    FROM "Album" AS a
        INNER JOIN "Track" AS t
            ON a."AlbumId" = t."AlbumId"
    GROUP BY a."ArtistId"
        , a."Title"
)
, js_albums AS
(
    SELECT row_to_json(albums) AS album_tracks
    FROM albums
)
, artist_albums AS
(
    SELECT a."Name" AS artist
        , array_agg(al.album_tracks) AS albums_tracks
    FROM "Artist" AS a
        INNER JOIN js_albums AS al
            ON a."ArtistId" = CAST(al.album_tracks->>'artist_id' AS INT)
    GROUP BY a."Name"
)
SELECT CAST(row_to_json(artist_albums) AS JSONB) AS artist_data
FROM artist_albums
;
```



Live with Chinook data



-- Select data from the view

```
SELECT *  
FROM v_artist_data  
;
```

	artist_data
1	{"artist": "The Cult", "albums_tracks": [{"artist_id": 139, "album_title": "Beyond Good And Evil", "album_tracks": ["War (The Process)", "The Saint", "Rise"]}
2	{"artist": "Boston Symphony Orchestra & Seiji Ozawa", "albums_tracks": [{"artist_id": 229, "album_title": "Carmina Burana", "album_tracks": ["Carmina"]}
3	{"artist": "Eugene Ormandy", "albums_tracks": [{"artist_id": 226, "album_title": "Respighi:Pines of Rome", "album_tracks": ["Pini Di Roma (Pinien Von Ro"]}
4	{"artist": "Barry Wordsworth & BBC Concert Orchestra", "albums_tracks": [{"artist_id": 224, "album_title": "The Last Night of the Proms", "album_tracks": ["The Last Night of the Proms"]}
5	{"artist": "Stevie Ray Vaughan & Double Trouble", "albums_tracks": [{"artist_id": 133, "album_title": "In Step", "album_tracks": ["The House Is Rockin'", "In Step"]}
6	{"artist": "Milton Nascimento", "albums_tracks": [{"artist_id": 42, "album_title": "Minas", "album_tracks": ["Minas", "Beijo Partido", "Ponta de Areia", "Tr"]}
7	{"artist": "Jorge Ben", "albums_tracks": [{"artist_id": 46, "album_title": "Jorge Ben Jor 25 Anos", "album_tracks": ["Engenho De Dentro", "Alcohol", "Mama"]}
8	{"artist": "Sergei Prokofiev & Yuri Temirkanov", "albums_tracks": [{"artist_id": 232, "album_title": "Prokofiev: Symphony No.1", "album_tracks": ["Sympl"]}

168 row(s) fetched - 21ms

Grid



Live with Chinook data



```
-- SELECT data from that VIEW, that does querying
SELECT jsonb_pretty(artist_data) pretty_artistdata
FROM v_artist_data
WHERE artist_data->>'artist' IN ('Miles Davis', 'AC/DC')
;
```

	T pretty_artistdata
1	{ "artist": "Miles Davis", "albums_tracks": [{ "artist_id": 68, "album_title": "The Essential Miles Davis [Disc 2]", "albu
2	{ "artist": "AC/DC", "albums_tracks": [{ "artist_id": 1, "album_title": "For Those About To Rock We Salute You", "albu


2 row(s) fetched - 18ms

Grid

✓ ✗ ↩ + (- ⏮ ⏪ ⏩ ⏭ 🔄 📄 ⚙



Live with Chinook data



```
-- SELECT some data that VIEW using JSON methods
SELECT jsonb_pretty(artist_data#>'{albums_tracks}') AS all_albums
      , jsonb_pretty(artist_data#>'{albums_tracks, 0}') AS tracks_0
      , artist_data#>'{albums_tracks, 0, album_title}' AS title
      , artist_data#>'{albums_tracks, 0, artist_id}' AS artist_id
      , artist_data->>'artist' AS artist
FROM v_artist_data
WHERE artist_data->'albums_tracks' @> '["album_title":"Miles Ahead"]'
;
```

	T all_albums	T tracks_0	? title	? artist_id	T artist
1	{ "artist_id": 68, "album_title": "The Essential Miles Davis [Disc 2]" }	{ "artist_id": 68, "album_title": "The Essential Miles Davis [Disc 2]" }	"The Essential Miles Davis [Disc 2]"	68	Miles Davis


```
{
  "artist_id": 68,
  "album_title": "The Essential Miles Davis [Disc 2]",
  "album_tracks":
  [
    "My Funny Valentine (Live)",
    "E.S.P.",
    "Nefertiti",
    "Petits Machins (Little Stuff)",
  ]
}
```

1 row(s) fetched - 19ms

Grid



Live with Chinook data



-- Array to records

```
SELECT jsonb_array_elements(artist_data#>'{albums_tracks}')->>'artist_id' AS artist_id  
      , artist_data->>'artist' AS artist  
      , jsonb_array_elements(artist_data#>'{albums_tracks}')->>'album_title' AS ablum_title  
      , jsonb_array_elements(jsonb_array_elements(artist_data#>'{albums_tracks}')#>'{album_tracks}') AS song_titles  
FROM v_artist_data  
WHERE artist_data->'albums_tracks' @> '[{"artist_id":139}]'  
ORDER BY 3, 4;
```

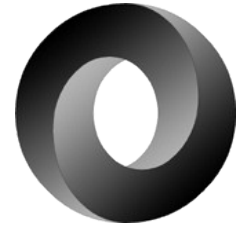
	T artist_id 📌	T artist 📌	T ablum_title 📌	? song_titles 📌
1	139	The Cult	Beyond Good And Evil	"American Gothic"
2	139	The Cult	Beyond Good And Evil	"Breathe"
3	139	The Cult	Beyond Good And Evil	"Heart Of Soul"
4	139	The Cult	Beyond Good And Evil	"Lil' Evil"
5	139	The Cult	Beyond Good And Evil	"Rain"
6	139	The Cult	Beyond Good And Evil	"Resurrection Joe"
7	139	The Cult	Beyond Good And Evil	"Rise"
8	139	The Cult	Beyond Good And Evil	"Shape The Sky"

30 row(s) fetched - 17ms

Grid



Live with Chinook data



```
-- Create a function, which will be used for UPDATE on the view v_artrist_data
CREATE OR REPLACE FUNCTION trigger_v_artist_data_update()
    RETURNS trigger AS
$BODY$
    -- Data variables
    DECLARE rec                RECORD;
    -- Error variables
    DECLARE v_state            TEXT;
    DECLARE v_msg              TEXT;
    DECLARE v_detail           TEXT;
    DECLARE v_hint             TEXT;
    DECLARE v_context          TEXT;
BEGIN
    -- Update table Artist
    IF OLD.artist_data->>'artist' <> NEW.artist_data->>'artist' THEN
        UPDATE "Artist"
        SET "Name" = NEW.artist_data->>'artist'
        WHERE "ArtistId" = artist_data#>'{albums_tracks, 0, artist_id}';
    END IF;

    -- Update table Album in a foreach
    -- Update table Track in a foreach

    RETURN NEW;
    EXCEPTION WHEN unique_violation THEN
        RAISE NOTICE 'Sorry, but the something went wrong while trying to update artist data';
    RETURN OLD;
    WHEN others THEN
        GET STACKED DIAGNOSTICS
            v_state = RETURNED_SQLSTATE,
            v_msg = MESSAGE_TEXT,
            v_detail = PG_EXCEPTION_DETAIL,
            v_hint = PG_EXCEPTION_HINT,
            v_context = PG_EXCEPTION_CONTEXT;
        RAISE NOTICE '%', v_msg;
        RETURN OLD;
END;
$BODY$
LANGUAGE plpgsql;
```



Live with Chinook data



```
-- The trigger will be fired instead of an UPDATE statement to save data
CREATE TRIGGER v_artist_data_instead_update INSTEAD OF UPDATE
  ON v_artist_data
  FOR EACH ROW
  EXECUTE PROCEDURE trigger_v_artist_data_update();
```



Live Amazon reviews



```
-- Create a table for JSON data with 1998 Amazon reviews  
CREATE TABLE reviews(review_jsonb jsonb);
```

Name	Value
Query	CREATE TABLE reviews(review_jsonb jsonb)
Updated Rows	0
1 row(s) fetched - 32ms	



Live Amazon reviews



```
-- Import customer reviews from a file  
COPY reviews FROM '/var/tmp/customer_reviews_nested_1998.json'
```

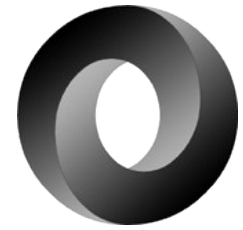
Name	Value
Query	-- Import customer reviews from a file COPY reviews FROM '/var/tmp/customer_reviews_nested_1998.json'
Updated Rows	0
1 row(s) fetched - 10730ms	

```
-- Maintenance the filled table  
VACUUM ANALYZE reviews;
```

Name	Value
Query	-- Maintenance the filled table VACUUM ANALYZE reviews
Updated Rows	0
1 row(s) fetched - 2371ms	



Live Amazon reviews

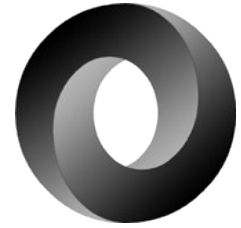


-- There should be 589.859 records imported into the table
SELECT count(*) FROM reviews;

	count
1	589.859
1 row(s) fetched - 104ms	



Live Amazon reviews



-- Select data with JSON

SELECT

review_jsonb#>> '{product,title}' **AS** title

, **avg**((review_jsonb#>> '{review,rating}')::int) **AS** average_rating

FROM reviews

WHERE review_jsonb@> '{"product": {"category": "Sheet Music & Scores"}}'

GROUP BY 1

ORDER BY 2 **DESC**

;

Without an Index: 248ms

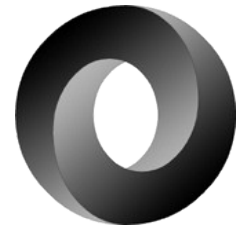
	title	average_rating
1	Complete Works for Solo Keyboard	5
2	The Magic Flute (Die Zauberflote in Full Score)	5
3	Requiem in Full Score	5
4	The Four Seasons and Other Violin Concertos in Full Score	5
5	Symphony No. 3 (Dover Miniature Scores)	5

12 row(s) fetched - 248ms

Grid



Live Amazon reviews




-- Create a GIN index

CREATE INDEX review_review_jsonb **ON** reviews **USING** GIN (review_jsonb);

Name	Value
Query	-- Create a GIN index CREATE INDEX review_review_jsonb ON reviews USING GIN (review_jsonb)
Updated Rows	0
1 row(s) fetched - 21079ms	



Live Amazon reviews



-- Select data with JSON

SELECT

review_jsonb#>> '{product,title}' **AS** title

, **avg**((review_jsonb#>> '{review,rating}')::int) **AS** average_rating

FROM reviews

WHERE review_jsonb@> '{"product": {"category": "Sheet Music & Scores"}}'

GROUP BY 1

ORDER BY 2 **DESC**

;

The same query as before with the previously created GIN Index: 7ms

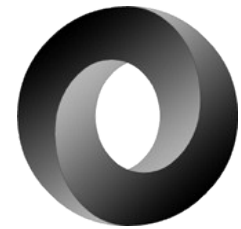
	title	average_rating
1	Complete Works for Solo Keyboard	5
2	The Magic Flute (Die Zauberflote in Full Score)	5
3	Requiem in Full Score	5
4	The Four Seasons and Other Violin Concertos in Full Score	5
5	Symphony No. 3 (Dover Miniature Scores)	5

12 row(s) fetched - 7ms

Grid



Live Amazon reviews



```
-- SELECT some statistics from the JSON data
SELECT review_jsonb#>>'{product,category}' AS category
      , avg((review_jsonb#>>'{review,rating}')::int) AS average_rating
      , count((review_jsonb#>>'{review,rating}')::int) AS count_rating
FROM reviews
GROUP BY 1
;
```

Without an Index: 9747ms

	category	average_rating	count_rating
1		4,487	1.521
2	Accessories	4,703	37
3	Action & Adventure	4,261	3.938
4	African American Cinema	4,694	36
5	Alternative Rock	4,522	15.508

84 row(s) fetched - 9747ms

Grid



Live Amazon reviews



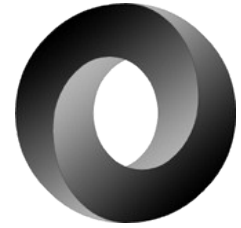
-- Create a B-Tree index on a JSON expression

CREATE INDEX reviews_product_category **ON** reviews ((review_jsonb#>>'{product,category}'));

Name	Value
Query	-- Create a B-Tree index on a JSON expression CREATE INDEX reviews_product_category ON reviews ((review_jsonb#>>'{product,category}'))
Updated Rows	0
1 row(s) fetched - 11875ms	



Live Amazon reviews



```
-- SELECT some statistics from the JSON data
SELECT review_jsonb#>>'{product,category}' AS category
      , avg((review_jsonb#>>'{review,rating}')::int) AS average_rating
      , count((review_jsonb#>>'{review,rating}')::int) AS count_rating
FROM reviews
GROUP BY 1
;
```

The same query as before with the previously created BTREE Index: 1605ms

	category	average_rating	count_rating
1		4,487	1.521
2	Accessories	4,703	37
3	Action & Adventure	4,261	3.938
4	African American Cinema	4,694	36
5	Alternative Rock	4,522	15.508

84 row(s) fetched - 1605ms

Grid



User Group / Conferences

PostgreSQL User Group Berlin

PGDay FOSDEM

Brussels, January 29 2016

FOSDEM Dev Room

Brussels, January 31 2016



Link List

Extended JSONB functions for 9.5

Extended JSONB functions as extension for 9.4
from [PGXN \(jsonbx\)](#)

Slide and source on Github:

<https://github.com/sjstoelting/talks>



JSON by the other Elephant



This document by [Stefanie Janine Stölting](#) is covered by the [Creative Commons Attribution 4.0 International](#)